

# Using DB2 LUW with DSX (Data Science Experience) to solve practical machine learning problems

**Mark Ryan**

*IBM Canada*

Session code: G02

Mon, Apr 30, 2018 (02:00 PM - 03:00 PM)

Db2 LUW

## Agenda

- Introduce Data Science Experience (DSX)
- Show steps to set up Db2 as a data source for DSX
- Show steps to complete a working example in DSX that applies deep learning to solve a practical problem
- Next steps & resources

## Who will find this session useful

- Familiar with Db2
- Not a machine learning guru
- Wants to take advantage of DSX / Watson Studio
- Looking for a straightforward way to apply machine learning techniques to everyday operational problems involving Db2 family data sources

## Why machine learning matters

- Artificial intelligence is yielding results
- A shift is happening “from programming computers to showing computers”
- ML is becoming to our industry what calculus is to mathematics
- Good news: data analysis is the key prerequisite for ML

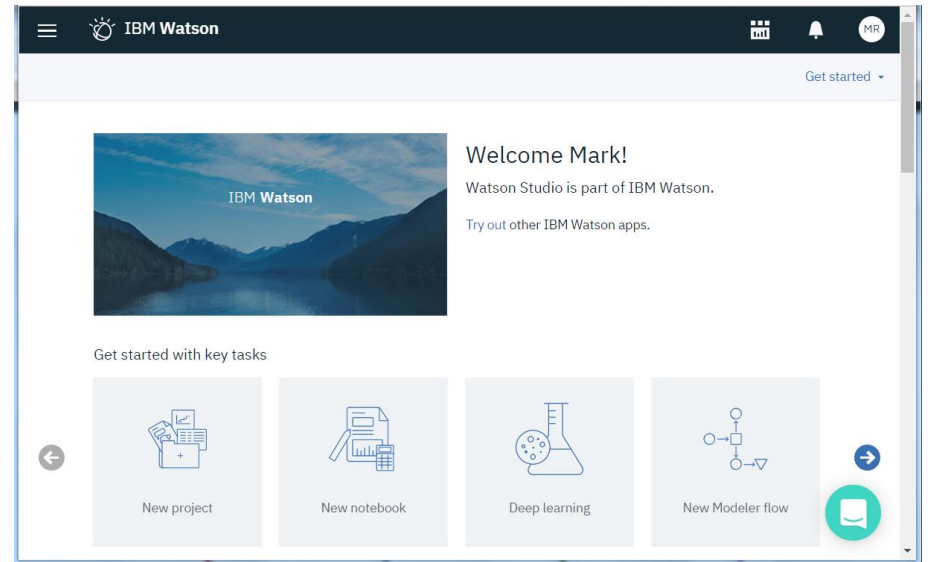
## Why DSX matters

- Machine learning is undergoing a democratization
- DSX provides a self-contained, end-to-end, accessible way to apply machine learning
- You don't need to be a Data Scientist to get a lot out of DSX!
- Many data sources available to tap into via DSX, including Db2
- The public web instantiation of DSX is now part of Watson Studio



# What is Data Science Experience (DSX)?

- DSX / Watson Studio includes tutorials, tools, articles, data sets and working code examples in the form of notebooks
- These capabilities also available in:
  - DSX Local
  - DSX Desktop
  - Integrated Analytics System
  - IBM Cloud Private for Data



# What are notebooks?

- Notebooks are working code examples that incorporate both human-readable documentation and executable code (e.g. Scala, Python)
- Sample Notebooks in DSX range from end-to-end scenarios (e.g. classifying tumours) to specific actions ( e.g. access Db2 from Scala)

## Apply linear regression on the prepared data

Apply multiple linear regression on the prepared data to get the prediction model.

```
# Linear Multiple Regression
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(v10X, v10Y)
print("Coef: " + str(regressor.coef_))
print("Intercept: " + str(regressor.intercept_))
```

```
Coef: [ 0.00878584  0.02903981 -0.00335873]
Intercept: 0.0604897840556
```

# Steps to set up Db2 as a data source for DSX

- Create required items:

Object	Where created	Credentials
Secure Gateway	IBM Cloud	Local host & port (for definition of the Destination with the Secure Gateway)
Client	Db2 system	NA
API call in notebook to access Db2	DSX	Local ID & password; Cloud host & port from Secure Gateway Destination properties

- Update Python code in notebook to refer to the data source with the appropriate credentials



## Steps to set up a Db2 data source for DSX – (1 | 3)

- Create Secure Gateway – details [here](#); summary of steps below:
  1. Log into [IBM Cloud](#) – from Dashboard select Create Service
  2. Search on Secure Gateway and select Free for pricing plan and Create
  3. Follow steps to create a Client
  4. Follow steps to create a Destination using the credentials for your Db2 system
  5. You will get the following credentials needed in subsequent steps:
    - Gateway ID
    - Cloud host
    - Cloud port

## Steps to set up a Db2 data source for DSX – (2|3)

- Set up Client on Db2 system
  1. Log into [IBM Cloud](#) – from Secure Gateway download appropriate client (Docker a good bet)
  2. Install Client on Db2 system
  3. In Docker environment:
    - start client using Gateway ID from Secure Gateway Settings:  
`Docker run -it ibmcom/secure-gateway-client pqDl217spyq_prod_ng`
    - run acl command acl allow host:port for Db2 system:  
`acl allow 10.0.9.93:3700`

## Steps to set up a Db2 data source for DSX – (3 | 3)

- Update notebook in DSX to tie everything together:
  1. Specify correct credentials:
    - Cloud host & port from Secure Gateway Destination Object settings
    - User and password from Db2 system
  2. You have a several options for the Python object to connect with your Db2 table – example below uses IdaDatabase object

```
dsn_uid = " "; # e.g. dbi04434
dsn_pwd = " "; # e.g. xxxx
dsn_hostname = ".integration.ibmcloud.com" # e.g. awh-yp-small03.services.dal.bluemix.net
dsn_port = "16833" # e.g. 50001
dsn_database = "SAMPLE" # e.g. BLUDB
```

### Create the database connection

The following code snippet creates a connection string `connection_string` and uses the `connection_string` to create a Db2 connection object

```
connection_string='jdbc:db2://' + dsn_hostname + ':' + dsn_port + '/' + dsn_database + ';user=' + dsn_uid + ';password=' + dsn_pwd + ';'
idadb=IdaDataBase(dsn=connection_string)
```

## Using Deep Learning in DSX to solve a practical problem

- A key aspect of a Db2 ticket is **Time to relief** (TTR) – the time taken to provide the customer with a way to address their primary symptom
- It would be very helpful to be able to predict as soon as a ticket is opened whether TTR will be more or less than 24 hours
- Goal: apply deep learning in DSX to predict when a ticket is first opened whether TTR will be greater than or less than 24 hours

## Why this example is relevant

- Deep learning is the machine learning approach that is currently producing the most ground-breaking results
- Traditionally, deep learning has been applied mostly to unstructured data (e.g. Images, audio)
- This example is a simple illustration of:
  - Relatively new approaches to applying deep learning to structured data
  - How to apply deep learning to data in Db2 using DSX

## Steps in this project (1 of 4)

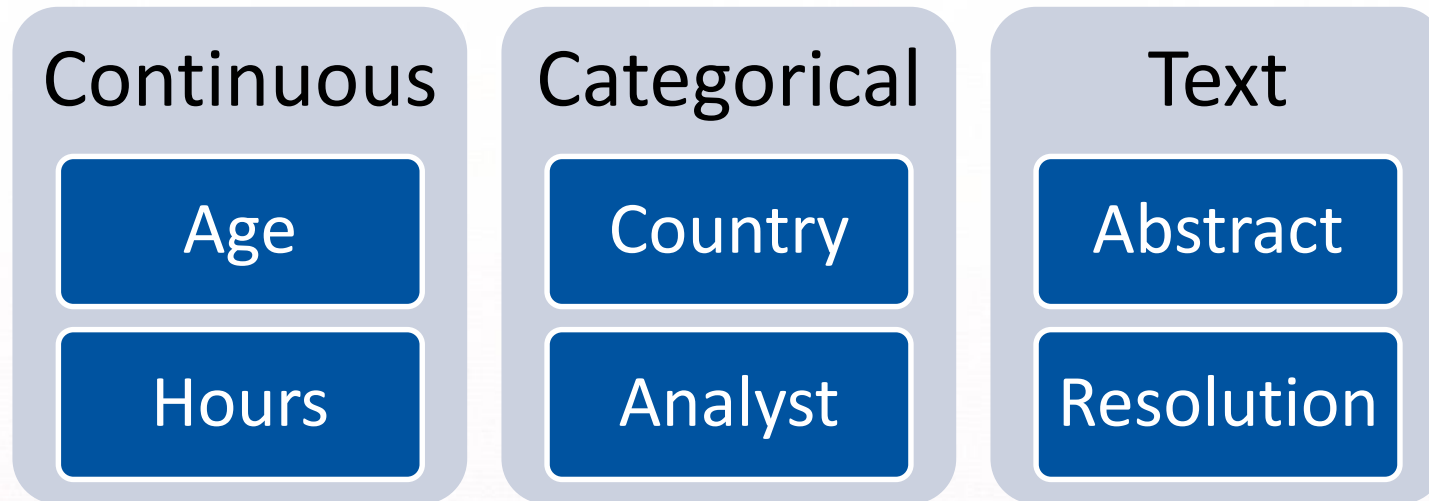
- Research approaches to deep learning with structured data:
  - [Great article](#) on structured deep learning
  - [Kaggle competition](#) with great examples of applying DL to structured data
  - [Fantastic course](#) covering examples of deep learning on structured data
- Select the framework:
  - Notebooks in DSX as a development environment
  - Keras as a deep learning framework:
    - Open
    - Flexible enough for this application without being overwhelmingly complex (contrast with Tensor Flow, the framework on which Keras abstracts)
    - Widely used – large community; many questions already answered



## Steps in this project (2 of 4)

### 3. Select features (columns) to feed into the model:

- Pick only from subset of data that is available when the ticket is first opened
- Avoid dependent features
- Selected features covered three classes of data:



## Steps in this project (3 of 4)


- Iteratively run model:
  - Tune hyperparameters (e.g. Learning rate, parameters to control overfitting, batch size)
  - Experiment with model structure (e.g. Add additional dense layers prior to output, try different optimization and activation functions)
- Goal: best possible validation accuracy:
  - Accuracy of the model's prediction of TTR <> 24 hours vs. Actual TTR for the validation set

## Steps in this project (4 of 4)

- Run experiments on input features:
  - Experiment with subsets of the available data by feature class:
    - Text feature only (Abstract)
    - Categorical/continuous features only
    - Combination of all three classes of features
  - Experiment with calculated features:
    - Include time since GA of release as a calculated continuous feature
    - Exclude zero-day TTR
    - Distort ticket severity feature to highlight extremes (Sev 1 and Sev 4)
    - Include day of week, month, year, day as calculated continuous features
  - Experiment with scope of the input data set:
    - Db2 only
    - Db2 + other Analytics products

## Lessons learned

1. Keras in DSX a solid deep learning platform
2. Tuning hyperparameters can be a distraction:
  - Fundamental problems with accuracy were solved by correcting input data & model problems
  - Tuning hyperparameters helped stability but not accuracy
3. Avoiding overfitting is key
4. More data = better accuracy



## Demo

```
BATCH_SIZE = 2000
epochs = 20
print("LR ", learning_rate)
print("dropout ", dropout_rate)
print("L2 lambda ", l2_lambda)
print("batch size ", BATCH_SIZE)
print("epochs ", epochs)

model = get_model()
model.fit(X_train, dtrain.target, epochs=epochs, batch_size=BATCH_SIZE
        , validation_data=(X_valid, dvalid.target), verbose=1)

LR 0.001
dropout 0.003
L2 lambda 7.5
batch size 2000
epochs 20
Train on 105230 samples, validate on 26308 samples
Epoch 1/20
105230/105230 [=====] - 60s - loss: 269.3049 - acc: 0.7026 - val_loss: 184.2830 - val_acc: 0.6529
Epoch 2/20
105230/105230 [=====] - 54s - loss: 131.7222 - acc: 0.7649 - val_loss: 87.9015 - val_acc: 0.6950
Epoch 3/20
105230/105230 [=====] - 54s - loss: 61.5127 - acc: 0.7797 - val_loss: 39.9705 - val_acc: 0.6809
Epoch 4/20
105230/105230 [=====] - 54s - loss: 27.2616 - acc: 0.7904 - val_loss: 17.2707 - val_acc: 0.6927
```

## Next steps

- Apply the same framework to other problems:
  - Predicting tickets with low NPS
  - Predicting code changes that are likely to cause regressions
  - Predicting tickets that are likely to require code changes to resolve
  - Predicting tickets that are likely to result in Duty Manager calls
- Deployment of models in production



## Resources

- Andrew Ng [intro to ML](#)
- Andrew Ng [Deep Learning Curriculum](#)
- Jeremy Howard [Deep Learning](#)
- Competitions on [Kaggle](#)
- Code for the [example in this presentation](#)
- Fantastic end-to-end example of [using DSX with Db2 Warehouse on Cloud](#)
- Essential source for ideas & inspiration: [Towards Data Science](#)

**Mark Ryan**

**IBM**

[mryan@ca.ibm.com](mailto:mryan@ca.ibm.com)

Session code: G02

*Please fill out your session  
evaluation before leaving!*

